

Personal Information

Name: **Maarten Blom**

StudentID: **12877034**

Email: maarten.blom@student.uva.nl

Submitted on: **22.03.2024**

Data Context

The data used was found here: <https://www.kaggle.com/code/ekami66/detailed-exploratory-data-analysis-with-python>

This is a curated dataset. Not much cleaning will need to be performed.

The goal is to find provenance in slides through the papers like this:

slide1 is created from paper1

slide2 is created from paper2

paper1 refers to paper2

slide1 has provenance in slide 2

Data Description

Each entry contains a scientific paper and its slides, stored both in XML and PDF format.

Also XML files contain the structure and the content of the documents, but does not include the images.

XML is easy to work with, because it is very structured. The PDF files are harder to work with. I will try this approach later.

Imports

```
In [1]: #!pip install tqdm
```

```
In [2]: import re  
import os
```

```
import numpy as np
import pandas as pd
from tqdm import tqdm
import xml.etree.ElementTree as ET
import matplotlib.pyplot as plt
```

Data Loading

```
In [3]: df = pd.DataFrame(columns = ['Paper', 'Slides'])

for folder_num in tqdm(range(5001)):
    paper = np.nan
    slides = np.nan
    folder_path = os.path.join("dataset", str(folder_num))
    for file_name in os.listdir(folder_path):
        if file_name.endswith('.tei.xml'):
            paper = ET.parse(os.path.join(folder_path, file_name))
        if file_name.endswith('.clean_tika.xml'):
            slides = ET.parse(os.path.join(folder_path, file_name))

    df.loc[folder_num] = [paper, slides]
```

[illegible]

```
In [4]: print("amount of rows with missing .xml files: ", len(df[df.isna().any(axis=1)]), "  
print("these will be dropped")  
df = df.dropna()  
display(df)
```

amount of rows with missing .xml files: 17 out of 5001
these will be dropped

	Paper	Slides
0	<xml.etree.ElementTree.ElementTree object at 0...	<xml.etree.ElementTree.ElementTree object at 0...
1	<xml.etree.ElementTree.ElementTree object at 0...	<xml.etree.ElementTree.ElementTree object at 0...
2	<xml.etree.ElementTree.ElementTree object at 0...	<xml.etree.ElementTree.ElementTree object at 0...
3	<xml.etree.ElementTree.ElementTree object at 0...	<xml.etree.ElementTree.ElementTree object at 0...
4	<xml.etree.ElementTree.ElementTree object at 0...	<xml.etree.ElementTree.ElementTree object at 0...
...
4979	<xml.etree.ElementTree.ElementTree object at 0...	<xml.etree.ElementTree.ElementTree object at 0...
4980	<xml.etree.ElementTree.ElementTree object at 0...	<xml.etree.ElementTree.ElementTree object at 0...
4981	<xml.etree.ElementTree.ElementTree object at 0...	<xml.etree.ElementTree.ElementTree object at 0...
4982	<xml.etree.ElementTree.ElementTree object at 0...	<xml.etree.ElementTree.ElementTree object at 0...
4983	<xml.etree.ElementTree.ElementTree object at 0...	<xml.etree.ElementTree.ElementTree object at 0...

4984 rows × 2 columns

Raw Text extraction

Analysis 1: Paper text

First, get raw text.

```
In [5]: def get_content_from_paper(tree):
text = ""
for element in tree.iter():
    if element.tag == "{http://www.tei-c.org/ns/1.0}body":
        for sub_element in element.iter():
            if sub_element.tag is not None:
                text += sub_element.text or ""
return text

df.loc[:, "PaperContent"] = df.Paper.apply(get_content_from_paper)
```


Plot for length

```
In [6]: fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(12, 6))
fig.suptitle('Horizontally stacked subplots')
ax1.hist(df.PaperContent.apply(len))
ax2.hist(df.PaperContent.apply(len), range = (0, 3000))

print("Paper amount of characters Mean: ", df.PaperContent.apply(len).mean())
```

Paper amount of characters Mean: 28948.222110754414

Horizontally stacked subplots



An average text document page is ~3000 characters long. A few documents in the dataset are lower than that. They will need to be looked at individually.

```
In [7]: print("Amount of documents with a suspiciously low charcter count: ", len(df[df["Pa
```

Amount of documents with a suspiciously low charcter count: 32

Analysis 2: Slide text

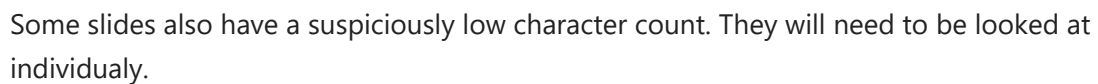
The same analysis can be done for slides

```
In [8]: def get_content_from_slides(tree):
    text = ""
    for element in tree.iter():
        if element.tag == "p":
            for sub_element in element.iter():
                if sub_element.tag is not None:
                    text += sub_element.text or ""
    return text

df.loc[:, "SlidesContent"] = df.Slides.apply(get_content_from_slides)
```

```
In [9]: fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(12, 6))
fig.suptitle('Horizontally stacked subplots')
```

Paper amount of characters Mean: 8708.395465489566



Analysis 3: Title and References

```
In [11]: #{http://www.tei-c.org/ns/1.0}

def get_title(tree):
    for element in tree.iter():
        if element.tag == "{http://www.tei-c.org/ns/1.0}titleStmt":
            for sub_element in element.iter():
                if sub_element.tag == "{http://www.tei-c.org/ns/1.0}title":
                    return sub_element.text
    return None

df.loc[:, "Title"] = df.Paper.apply(get_title)
```

In [12]: df

Out[12]:

	Paper	Slides	PaperCo
0	<xml.etree.ElementTree.ElementTree object at 0...	<xml.etree.ElementTree.ElementTree object at 0...	\nTheor Expo communication
1	<xml.etree.ElementTree.ElementTree object at 0...	<xml.etree.ElementTree.ElementTree object at 0...	\nIntroductionPe peer (P2P) file-sha
2	<xml.etree.ElementTree.ElementTree object at 0...	<xml.etree.ElementTree.ElementTree object at 0...	\nIntroduct ultimate achiev
3	<xml.etree.ElementTree.ElementTree object at 0...	<xml.etree.ElementTree.ElementTree object at 0...	\nINTRODUCTIO challenges of deve
4	<xml.etree.ElementTree.ElementTree object at 0...	<xml.etree.ElementTree.ElementTree object at 0...	\nIntroductionIr domains, converge
...	
4979	<xml.etree.ElementTree.ElementTree object at 0...	<xml.etree.ElementTree.ElementTree object at 0...	\nIntroductionPr Schumacher con
4980	<xml.etree.ElementTree.ElementTree object at 0...	<xml.etree.ElementTree.ElementTree object at 0...	\nIntroductionLan and therefore dat
4981	<xml.etree.ElementTree.ElementTree object at 0...	<xml.etree.ElementTree.ElementTree object at 0...	\nIntroduction auto-completion (C
4982	<xml.etree.ElementTree.ElementTree object at 0...	<xml.etree.ElementTree.ElementTree object at 0...	\nIntroductionCont or 'data-to-te
4983	<xml.etree.ElementTree.ElementTree object at 0...	<xml.etree.ElementTree.ElementTree object at 0...	\nIntroducti simplification (TS) i

4984 rows × 5 columns

Discussion

For a text-based approach, this is sufficient. After looking at the suspicious documents, I can train a NLP model on either the raw text or the entirety of the XML code of the slide and the paper.

I also want to include images found in the Paper and the Slides. These are not found in the XML files, but can be extracted from the PDF's. I've tried this, but haven't been able to filter out the layout elements and the actual figures.

After finding the corresponding paper to the slide, I can create a directed graph by linking the slides to the slides which its paper refers to. This is the goal that will be achieved.

In []: